

Mechanistic Interpretability Resource Guide

Getting Started

[Mechanistic Interpretability Quickstart Guide](#) [Neel Nanda, 2023]

[Concrete Steps to Get Started in Transformer Mechanistic Interpretability](#) [Neel Nanda, 2022]

[A Comprehensive Mechanistic Interpretability Explainer and Glossary](#) [Neel Nanda, 2022]

Explanatory Papers

[Zoom In: An Introduction to Circuits](#) (Olah, et.al. 2020)

[Curve Detectors](#) (Cammarata, et.al. 2020)

[Toy Models of Superposition](#) (Elhage, et.al. 2022)

Research Papers

[List of Mechanistic Interpretability Papers](#) [Neel Nanda]

[Multimodal Neurons in Artificial Neural Networks](#) (Goh, et.al. 2021)

[Towards Monosemanticity: Decomposing Language Models with Dictionary Learning](#) (Bricken, et.al. 2023)

[Scaling Monosemanticity: Extracting Interpretable Features from Claude 3 Sonnet](#) (Templeton and Conerly, et.al. 2024)

[Scaling and Evaluating Sparse Autoencoders](#) (Gao, et.al. 2024)

Code

[OpenAI Sparse Autoencoders GitHub](#) (sparse autoencoders trained on gpt2-small)

[Transformer Interpretability short course](#)

Videos

[Induction Head Circuits](#) [lecture by Chris Olah]

[Mechanistic Interpretability & Mathematics](#) [lecture by Neel Nanda]

Future Looking

[Interpretability Dreams](#) (Olah, et.al. 2023)

[200 Concrete Open Problems in Mechanistic Interpretability](#) [Neel Nanda, 2022]

[Open Questions](#) (from Toy Models of Superposition paper)

[Current Themes in MI Research](#) (article, 2022)